

STAT 135, Midterm exam, Spring 2016, H. Pitters

NAME (IN CAPS): _____
SID number and SECTION: _____

Show your work or provide a brief explanation for all answers. This quiz is closed books. You are allowed to use the notes you took during class, a calculator and extra scratch paper. Answers should be simplified as much as possible. Good luck!

1	2	3	4	5	Σ
---	---	---	---	---	----------

Question 1. [16] Provide a concise answer to each of the following questions.

- (1) Explain in your own words how the maximum likelihood method works. Is this method intuitively plausible? [2]

A: Given observations x_1, \dots, x_n from some random variables X_1, \dots, X_n with joint distribution \mathbb{P}_θ that depends on a parameter θ , the likelihood of θ is

$$\text{lik}(\theta) := f(x_1, \dots, x_n | \theta),$$

where $f(x_1, \dots, x_n | \theta)$ is the joint density of X_1, \dots, X_n under \mathbb{P}_θ . The MLE for θ is the value $\hat{\theta}$ that maximizes $\text{lik}(\theta)$. This is a plausible estimator for θ , since it corresponds to the one model $\mathbb{P}_{\hat{\theta}}$ under which the observations are most likely to be observed.

- (2) Explain what it means to have a 95% confidence interval for some parameter μ . [2]

A: A 95% confidence interval is a random interval $[a(X), b(X)]$ (usually constructed from the sample $X = (X_1, \dots, X_n)$). If one were to repeatedly and independently draw samples of size n and construct a 95% confidence interval for each of them, then about 95% of these intervals will contain the parameter θ .

- (3) Based on a sample of size n , drawn from an unknown population, we are told that $[2.7, 5.2]$ is a 90% confidence interval for some unknown parameter μ . Does this mean the probability for μ to lie in $[2.7, 5.2]$ is 90%? Explain your answer. [2]

A: No. The parameter μ is deterministic, and so is the interval $[2.7, 5.2]$. It either contains μ or it does not.

- (4) Suppose $\hat{\theta}$ is an estimator for some unknown parameter θ . What does it mean for $\hat{\theta}$ to be unbiased? [2]

A: It means that the mean value of $\hat{\theta}$ is θ , i.e. $\mathbb{E}\hat{\theta} = \theta$.

- (5) Suppose T is a sufficient statistic for a parameter θ . [2]

(a) Explain in your own words what sufficiency means.

A: As far as inference (e.g. estimation) of θ is concerned, the statistic T contains all necessary information. Any information about the sample X_1, \dots, X_n other than $T(X_1, \dots, X_n)$ does not improve the inference of θ .

(b) Why is the maximum likelihood estimator (MLE) for θ a function of T ?

A: From the factorization theorem we have that

$$\text{lik}(\theta) = f(x_1, \dots, x_n | \theta) = g(T(x_1, \dots, x_n), \theta)h(x).$$

The MLE $\hat{\theta}$ is the value that maximizes $\text{lik}(\theta)$, or equivalently $g(T(x_1, \dots, x_n), \theta)$. However, the latter depends on x_1, \dots, x_n only via $T(x_1, \dots, x_n)$, hence $\hat{\theta}$ is a function of T .

(6) Explain the difference between the Central Limit Theorem (CLT) and the Law of Large Numbers (LLN). Assuming both theorems hold, can you derive one from the other? [2]

A: Consider the average of the first n items

$$\bar{X}_n := \frac{1}{n} \sum_{i=1}^n X_i$$

of a sequence X_1, X_2, \dots of independent and identically distributed random variables with mean $\mu := \mathbb{E}X_1$ and variance $\sigma^2 := \text{Var}(X_1)$. The Law of Large Numbers states that as n grows without bounds \bar{X}_n converges towards μ .

Even if n is large, the average $\frac{1}{n} \sum_{i=1}^n X_i$ is going to fluctuate randomly about μ . The Central Limit Theorem states that these fluctuations have a Gaussian distribution, namely the standardized average

$$\frac{\frac{1}{n} \sum_{i=1}^n X_i - \mu}{\sigma / \sqrt{n}}$$

has approximately a standard normal distribution $\mathcal{N}(0, 1)$.

Assuming both theorems hold, for any $\epsilon > 0$ the CLT implies

$$\mathbb{P} \left\{ \left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right| \geq \epsilon \right\} = \mathbb{P} \left\{ \left| \frac{\frac{1}{n} \sum_{i=1}^n X_i - \mu}{\sigma / \sqrt{n}} \right| \geq \frac{\sqrt{n}}{\sigma} \epsilon \right\} \approx 1 - \Phi\left(\frac{\sqrt{n}}{\sigma} \epsilon\right) \rightarrow 0$$

as $n \rightarrow \infty$. Alternatively, Chebyshev's inequality implies for any $\epsilon > 0$

$$\mathbb{P} \left\{ \left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right| \geq \epsilon \right\} \leq \frac{\sigma^2}{\epsilon^2 n} \rightarrow 0$$

as $n \rightarrow \infty$. Both statements show the so-called weak Law of Large Numbers.

- (7) For random variables X_1, \dots, X_n is the statement

$$\mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n X_i \right] = \frac{1}{n} \sum_{i=1}^n X_i$$

correct? Explain. [2]

A: The statement is false, since the left hand side is a (deterministic) number, whereas the right hand side is a random variable. For instance, with $n = 1$ the statement reads $\mathbb{E}X_1 = X_1$.

- (8) You are advising a political candidate on opinion polls. Two organizations have submitted bids to do a survey to estimate the extent of his support. Both will use nationwide probability samples of 2,500 people, and both charge the same price. One organization only states that it has a 95% chance to get the estimate right to within 3%. The other absolutely guarantees that its estimate will be right to within 3%. Other things being equal, which bid should be accepted, and why? [2]

A: Simply because of chance error there is no way that an estimate constructed from a sample (whose size is a small fraction of the total population) should be guaranteed to be right within 3%. The modest bid should therefore be accepted.

Question 2. [4] In a simple random sample of 100 graduate students from a certain college, 48 were earning \$70,000 a year or more. Estimate the percentage of all graduates of that college earning \$70,000 a year or more. [2] Provide an (estimated) standard error (ignoring finite population correction). [2]

A: An (unbiased) estimator \hat{p} for the percentage p of graduates earning \$70,000 a year or more is the fraction of graduates earning \$70,000 a year or more in the sample, i.e. $\hat{p} = 48\%$. An estimated standard error for this estimator is

$$\sigma_{\hat{p}} = \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \approx \frac{1}{20} = 0.05.$$

Question 3. [4] A sample X_1, \dots, X_n of n independent draws is taken with common distribution $\text{Pn}(\alpha)$. We want to perform Bayesian inference, where the parameter α is modeled as random with prior distribution an exponential distribution with parameter $\lambda > 0$.

- (1) Compute the joint distribution of X_1, \dots, X_n and α . [1]

A: Denoting by $\pi(\alpha) = \lambda e^{-\lambda\alpha} \mathbf{1}_{[0,\infty)}(\alpha)$ the density of the prior distribution of α , we have

$$\begin{aligned} f(x_1, \dots, x_n, \alpha) &= f(x_1, \dots, x_n | \alpha) \pi(\alpha) = \prod_{i=1}^n e^{-\alpha} \frac{\alpha^{x_i}}{x_i!} \lambda e^{-\lambda\alpha} \mathbf{1}_{[0,\infty)}(\alpha) \\ &= e^{-\alpha(n+\lambda)} \frac{\alpha^{\sum_i x_i}}{\prod_i x_i!} \lambda \mathbf{1}_{[0,\infty)}(\alpha). \end{aligned}$$

- (2) Compute the posterior distribution $f(\alpha | x_1, \dots, x_n)$ of α given the observations $X_1 = x_1, \dots, X_n = x_n$. (Hint: There is no need to compute the marginal distribution of X_1, \dots, X_n .) [3]

A: Denoting by $m(x_1, \dots, x_n)$ the probability mass function of the observations, we find

$$f(\alpha | x_1, \dots, x_n) = \frac{f(x_1, \dots, x_n, \alpha)}{m(x_1, \dots, x_n)} \propto e^{-\alpha(n+\lambda)} \alpha^{\sum_i x_i} \mathbf{1}_{[0,\infty)}(\alpha).$$

Consequently, the posterior distribution of α is a gamma($\sum_i x_i + 1, n + \lambda$) distribution.

Question 4. [16] In n independent experiments physicists measure the (unknown) mass μ of an object recording the result of the first experiment as X_1 , the result of the second experiment as X_2 , etc. up until X_n , the result of the n th experiment. The experimenters assume that measurements vary randomly about the true mass μ according to a normal distribution with variance 1, i.e. for each $1 \leq i \leq n$ they assume

$$X_i \sim \mathcal{N}(\mu, 1).$$

- (1) Find the method of moments estimator for μ . [2]

A: Population mean: $\mu_1 := \mathbb{E}X_1 = \mu$, hence $\hat{\mu} = \hat{\mu}_1 = \bar{X}_n$ is the method of moments estimator for μ .

- (2) Given the data x_1, \dots, x_n of the experiments, write down the likelihood $\text{lik}(\mu)$ of μ . [2]

A:

$$\text{lik}(\mu) := \prod_{i=1}^n \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(x_i - \mu)^2} = (2\pi)^{-\frac{n}{2}} e^{-\frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2}.$$

- (3) Find an (simplified) expression for the log likelihood $l(\mu) := \log \text{lik}(\mu)$. [2]

A:

$$l(\mu) = -\frac{n}{2} \log(2\pi) - \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2.$$

- (4) Find the maximum likelihood estimator for μ . [2]

A: The derivative

$$l'(\mu) = \sum_{i=1}^n (x_i - \mu) = n\bar{x}_n - n\mu$$

of the log likelihood has root $\hat{\mu} = \bar{x}$. Since for $\mu > 0$ we have $l''(\mu) < 0$, $\hat{\mu} = \bar{x}_n$ maximizes the likelihood and is therefore the MLE for μ .

- (5) As an estimator for μ consider

$$\bar{X}_n := \frac{1}{n} \sum_{i=1}^n X_i.$$

- (a) Is \bar{X}_n an unbiased estimator? [2]

A: Yes, since $\mathbb{E}\bar{X}_n = \mathbb{E}X_i = \mu$.

- (b) Is \bar{X}_n a consistent estimator? [2]

A: Yes. The Law of Large Numbers implies the convergence (in probability)

$$\frac{1}{n} \sum_{i=1}^n X_i \rightarrow \mathbb{E}X_1 = \mu,$$

and hence \bar{X}_n is a consistent estimator for μ .

- (c) What is the standard error of \bar{X}_n ? [2]

A: The standard error is $\sqrt{\text{Var}(\bar{X}_n)} = \sqrt{\frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i)} = \frac{1}{\sqrt{n}}$.

- (d) What is the sampling distribution of \bar{X}_n ? [2]

A: \bar{X}_n is a convolution of n i.i.d. random variables each of which has distribution $\mathcal{N}(\mu/n, 1/n^2)$, hence $\bar{X}_n \sim \mathcal{N}(\mu, 1/n)$.

Question 5. [2] Suppose $\hat{\theta}$ is an unbiased estimator for θ . Starting from the definition of the MSE, derive a simple expression for $\text{MSE}(\hat{\theta})$.

A:

$$\text{MSE}(\hat{\theta}) := \mathbb{E}(\hat{\theta} - \theta)^2 = \mathbb{E}\hat{\theta}^2 - 2\theta\mathbb{E}\hat{\theta} + \theta^2 = \mathbb{E}\hat{\theta}^2 - \theta^2 = \mathbb{E}\hat{\theta}^2 - (\mathbb{E}\hat{\theta})^2 = \text{Var}(\hat{\theta}),$$

since $\mathbb{E}\hat{\theta} = \theta$ by assumption.